

Assignment Project Exam Help

Linear Regression

<https://powcoder.com>

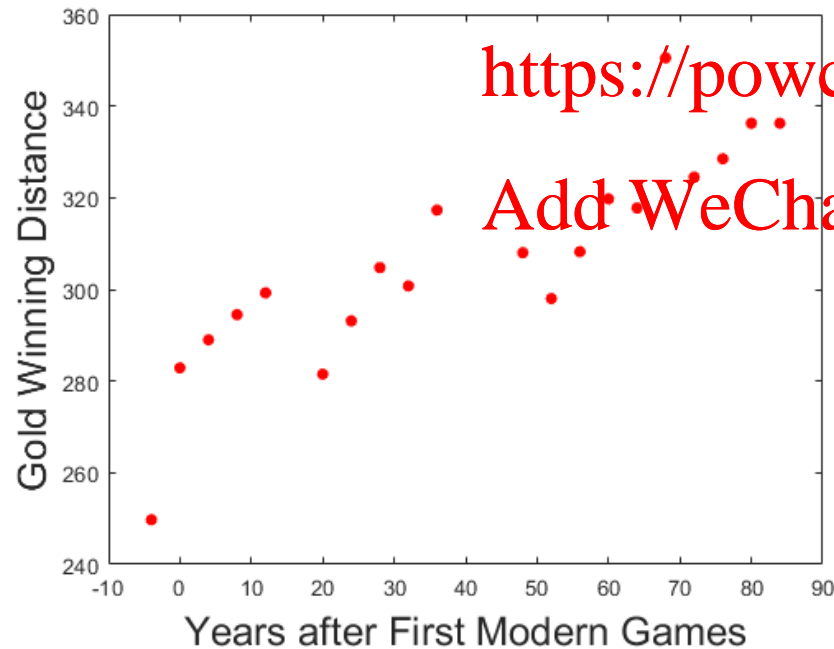
Add WeChat powcoder

Ata Kaban

Example of a regression problem

- Let's look at some fun data. Can we predict the long-jump Gold winning distance 125 years after the first modern Olympic games?

Assignment Project Exam Help



<https://powcoder.com>


Add WeChat powcoder

Data from: Rogers & Girolami. A First Course in Machine Learning. 2nd edition. Chapman & Hall/CRC, 2017

Example of a regression problem

- Let's look at more fun data. Can we predict people's weight from their height?

Height(cm)	Weight(kg)
150.00686	69.73347
151.64326	69.83261
155.54032	71.55736
158.80535	71.92875
161.17561	72.92118
...	...
175.15167	79.48533
182.32900	80.52182
191.11317	83.67998
193.21947	84.72086
202.68705	89.64049



Regression

- Regression means learning a function that captures the “trend” between input and output
<https://powcoder.com>
- We then use this function to predict target values for new inputs
[Add WeChat powder](#)

Univariate linear regression

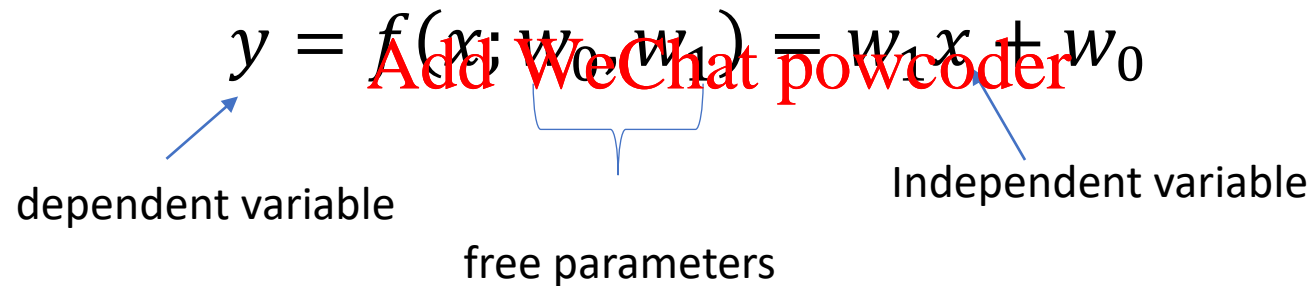
- Visually, there appears to be a trend
- A reasonable **model** seems to be the **class of linear functions (lines)**
- We have one input attribute (year) - hence the name **univariate**

$$y = f(x; w_0, w_1) = w_1 x + w_0$$

dependent variable

free parameters


Independent variable

The diagram shows the equation y = f(x; w_0, w_1) = w_1 x + w_0. A blue arrow points from the label 'dependent variable' to the variable 'y'. Another blue arrow points from the label 'Independent variable' to the variable 'x'. A blue bracket is placed under the parameters 'w_0, w_1', with a line pointing down to the label 'free parameters'.

- Any line is described by this equation by specifying values for w_1, w_0 .

Check your understanding

Height(cm)	Weight(kg)
150.00686	69.73347
151.64326	69.83261
155.54032	71.55730
158.80535	71.92875
161.17561	72.92118
...	
175.15167	79.48533
182.32900	80.52182
191.11317	83.67998
193.21947	84.72086
202.68705	89.64049



Suppose that from historical data someone already calculated the parameters of our linear model are $w_0 = 1.68$, $w_1 = 0.44$. A new person (James) has height $x=178\text{cm}$. Using our model, we can predict James' weight is $0.44 * 178 + 1.68 = 80\text{kg}$.

Play around with linear functions

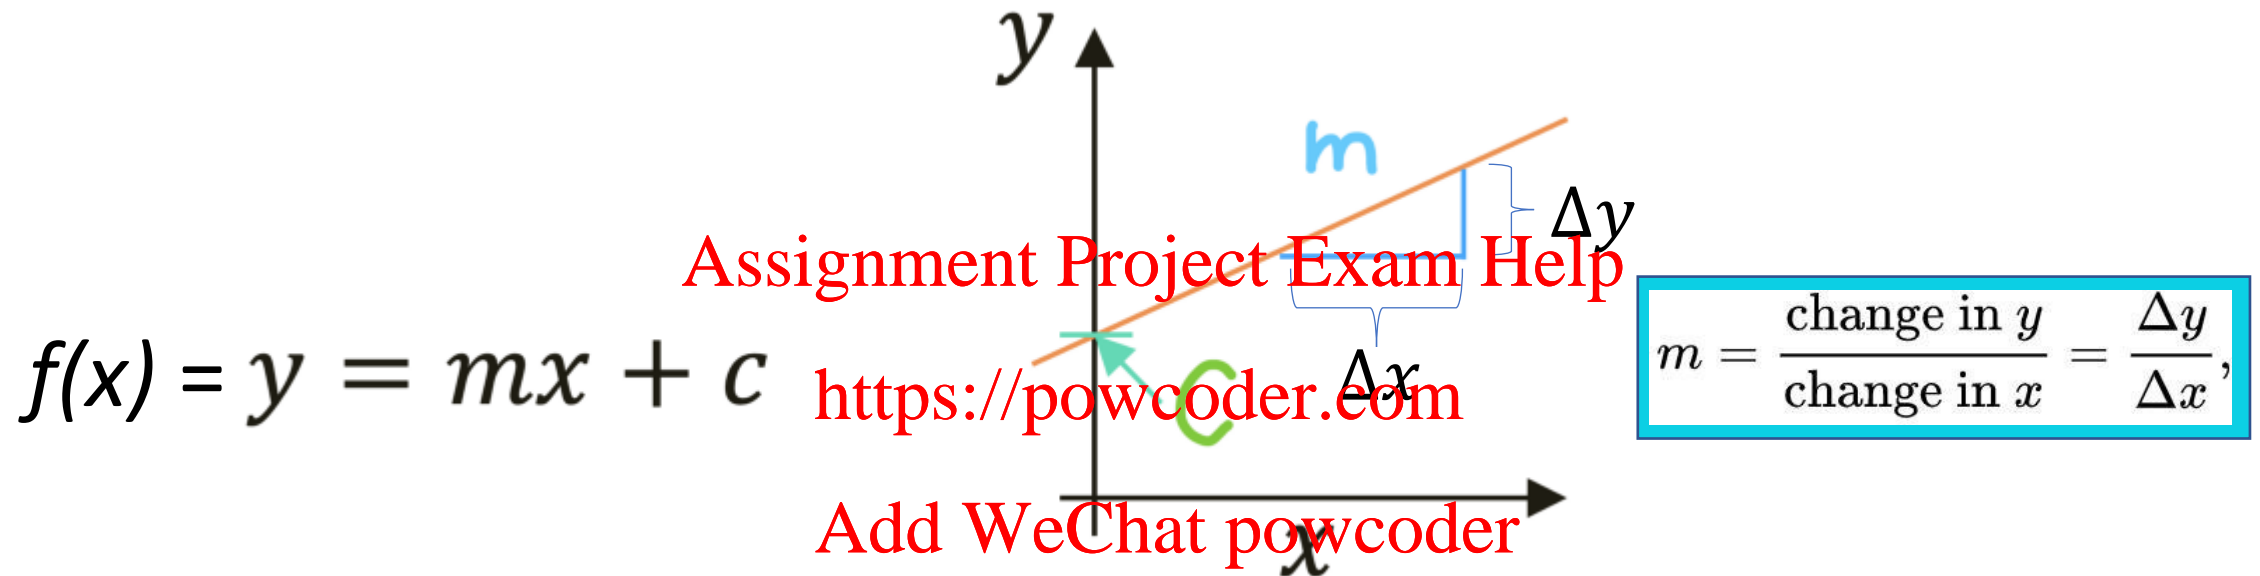
- Go to <https://www.desmos.com/calculator>
- Type: $y = w_1x + w_0$
- Plug in some values for the free parameters, or use the slider to see their effect
- *What is the role of the free parameters?*
 - w_1 is the slope of the line
 - w_0 is the intercept with the y-axis
- Fixing concrete numbers for these parameters gives you specific lines

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Equation of a straight line with slope m and intercept c .



This is why:

$$f(x + \Delta x) = m(x + \Delta x) + c = mx + m \cdot \Delta x + c = f(x) + m \cdot \Delta x$$
$$\Rightarrow m = \frac{f(x + \Delta x) - f(x)}{\Delta x} = \frac{\Delta y}{\Delta x}$$

Our goal: Find the “best” line



- Which is the “best” line? That captures the trend in the data.
- Determine the “best” values for w_1 , w_0 .

Loss functions (or cost functions)

- We need a criterion that, given the data, for any given line will tell us how bad is that line.

Assignment Project Exam Help

<https://powcoder.com>


- Such criterion is called a loss function. It is a function of the free parameters!

Add WeChat powcoder

Terminology

- Loss function = cost function = loss = cost = error function

We average the losses on all training examples



For each training example (point)
 $n = 1, \dots, N,$

The loss on the n -th point is the mismatch between the output of the model for this point $f(x^{(n)}; w_0, w_1)$ and the observed target $y^{(n)}$.

Average these losses.

Square loss (L2 loss)


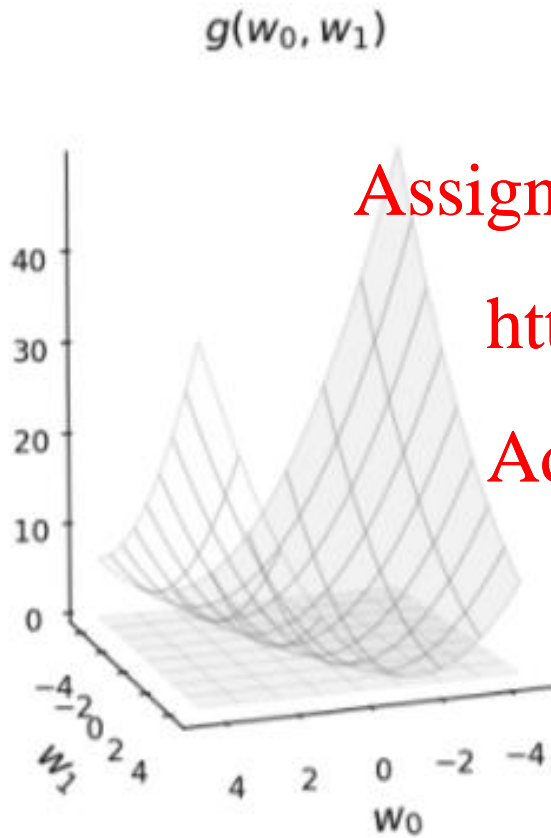
- The loss expresses an error, so it must be always non-negative
- Square loss is a sensible choice to measure mismatch for regression
- Mean Square Error (MSE)

$$g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N \underbrace{(f(x^{(n)}; w_0, w_1) - y^{(n)})^2}_{\text{loss for the n-th training example}}$$

loss for the n-th training example

and recall that, for any x , we have $f(x; w_0, w_1) = w_1 x + w_0$

Cost function depends on the free parameters



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Check your understanding

- Suppose a linear function with parameters $w_0 = 0.5$, $w_1 = 0.5$
- Compute the loss function value for this line at the training example: $(1, 3)$.

Assignment Project Exam Help

<https://powcoder.com>

- $f(x^{(1)}; 0.5, 0.5) = 0.5 * 1 + 0.5 = 1$ (output of the model)
- $y^{(1)} = 3$ (actual target)
- Square loss for this point: $(1-3)^2 = 4$.
- Cost = 4.

Add WeChat powcoder

Univariate linear regression – what we want to do

- Given training data

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(N)}, y^{(N)})$$

Assignment Project Exam Help

- Fit the model

<https://powcoder.com>

$$y = f(x; w_0, w_1) = w_1 x + w_0$$

Add WeChat powcoder

- By minimising the cost function

$$g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})^2$$

Univariate linear regression – what we want to do

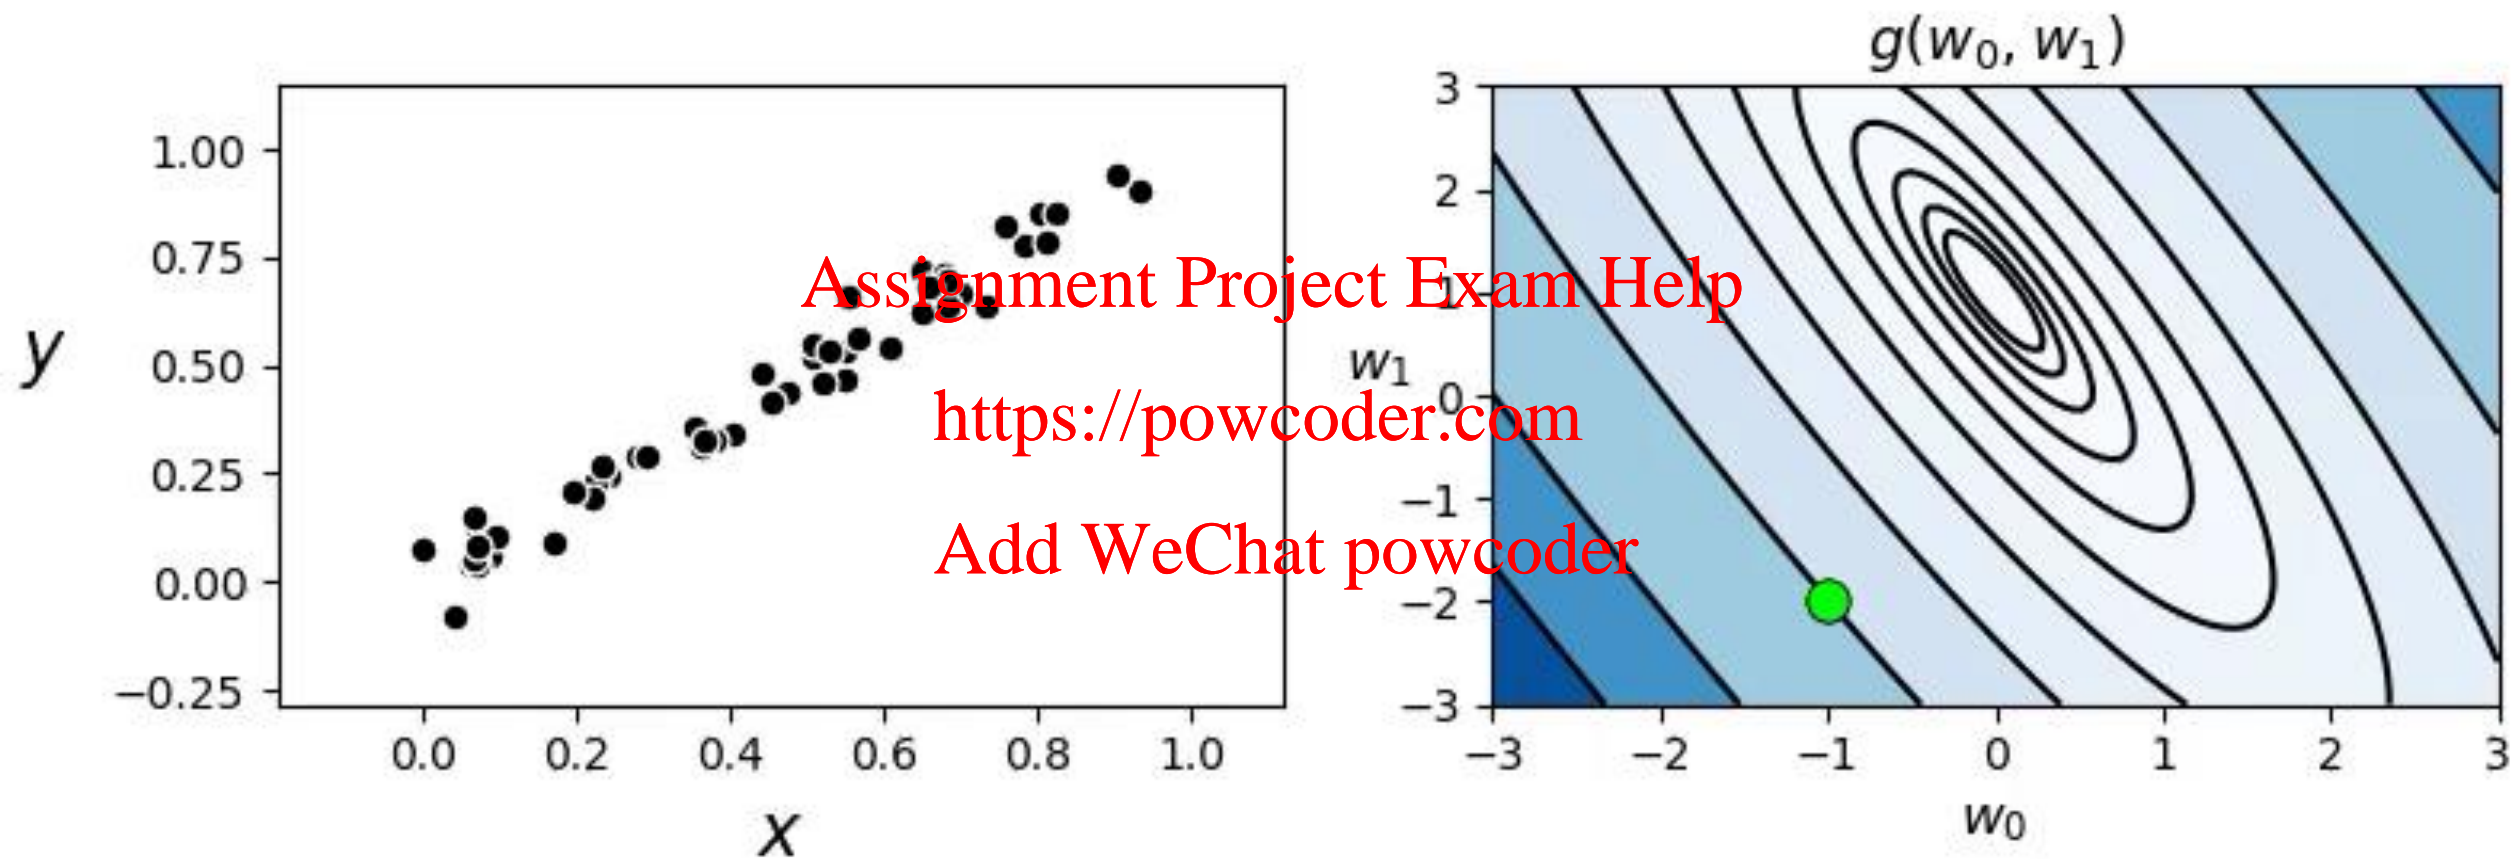
- Every combination of w_0 and w_1 has an associated cost.

Assignment Project Exam Help

- To find the ‘best fit’ we need to find values for w_0 and w_1 such that the cost is minimum.

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

Gradient Descent

<https://powcoder.com>

Add WeChat powcoder

Gradient Descent

- A general strategy to minimise cost functions.

Goal: Minimise cost function $g(w_0, w_1)$


Start at say $w_0 := 0, w_1 := 0$

Repeat until no change occurs

Update w_0, w_1 by taking

a small step in the direction of the steepest descent

Return w_0, w_1



Gradient descent – the general algorithm

- Goal: Minimise cost function $g(\mathbf{w})$, where $\mathbf{w} = (w_0, w_1, \dots)$

Input: $\alpha > 0$

Initialise \mathbf{w} // at 0 or some random value

Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \alpha \nabla g(\mathbf{w})$$

Return \mathbf{w}

step size

direction

α is called “learning rate” = “step size”, for instance 0.01


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

How to find the best direction?


- First, recall from calculus that the derivative of a function is the change in function value as the argument of the function changes by a minimal amount. <https://powcoder.com>
- The derivative evaluated at a given location gives us the slope of the tangent line at that point. Add WeChat powcoder
- The negative of the slope points towards the minimum point. Check!



$$\frac{\Delta y}{\Delta x} = \frac{f(x+\Delta x) - f(x)}{\Delta x}$$

$\Delta x \rightarrow 0$

Demo example for gradient descent algorithm



Gradient

- **Partial derivative** with respect to w_0 is $\frac{\delta g(w_0, w_1)}{\delta w_0}$. It means the derivative function of $g(w_0, w_1)$ when w_1 is treated as constant.
- **Partial derivative** with respect to w_1 is $\frac{\delta g(w_0, w_1)}{\delta w_1}$. It means the derivative function of $g(w_0, w_1)$ when w_0 is treated as constant.
- The vector of partial derivatives is called the **gradient**.

$$\nabla g(w) = \begin{pmatrix} \frac{\delta g(w_0, w_1)}{\delta w_0} \\ \frac{\delta g(w_0, w_1)}{\delta w_1} \end{pmatrix} \quad \text{where } w = \begin{pmatrix} w_0 \\ w_1 \end{pmatrix}$$

- The negative of the gradient evaluated at a location $(\widehat{w}_0, \widehat{w}_1)$ gives us the direction of the **steepest descent** from that location.
- We take a small step in that direction.

Assignment Project Exam Help
Gradient Descent applied to solving
Univariate Linear Regression
<https://powcoder.com>

Add WeChat powcoder

Computing the gradient for our L2 loss

- Recall the cost function $g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})^2$
- Using the chain rule, we have:*

$$\frac{\delta g(w_0, w_1)}{\delta w_0} = \frac{2}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})$$
$$\frac{\delta g(w_0, w_1)}{\delta w_1} = \frac{2}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)}) x^{(n)}$$

*For a very detailed explanations of all steps watch: <https://www.youtube.com/watch?v=sDv4f4s2SB8>



Algorithm for univariate linear regression using GD

- Goal: Minimise $g(w_0, w_1) = \frac{1}{N} \sum_{n=1}^N ((w_1 x^{(n)} + w_0) - y^{(n)})^2$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

```
Input:  $\alpha > 0$ , training set  $\{(x^{(n)}, y^{(n)}): n=1, \dots, N\}$   
Initialise  $w_0 := 0, w_1 := 0$   
Repeat  
    For  $n=1, \dots, N$  // more efficient to update after each data point  
         $w_0 := w_0 - \alpha \cdot ((w_1 x^{(n)} + w_0) - y^{(n)})$   
         $w_1 := w_1 - \alpha \cdot ((w_1 x^{(n)} + w_0) - y^{(n)}) x^{(n)}$   
Until change remains below a very small threshold  
Return  $w_0, w_1$ 
```

Effect of the learning rate



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Extensions & variants of regression problems

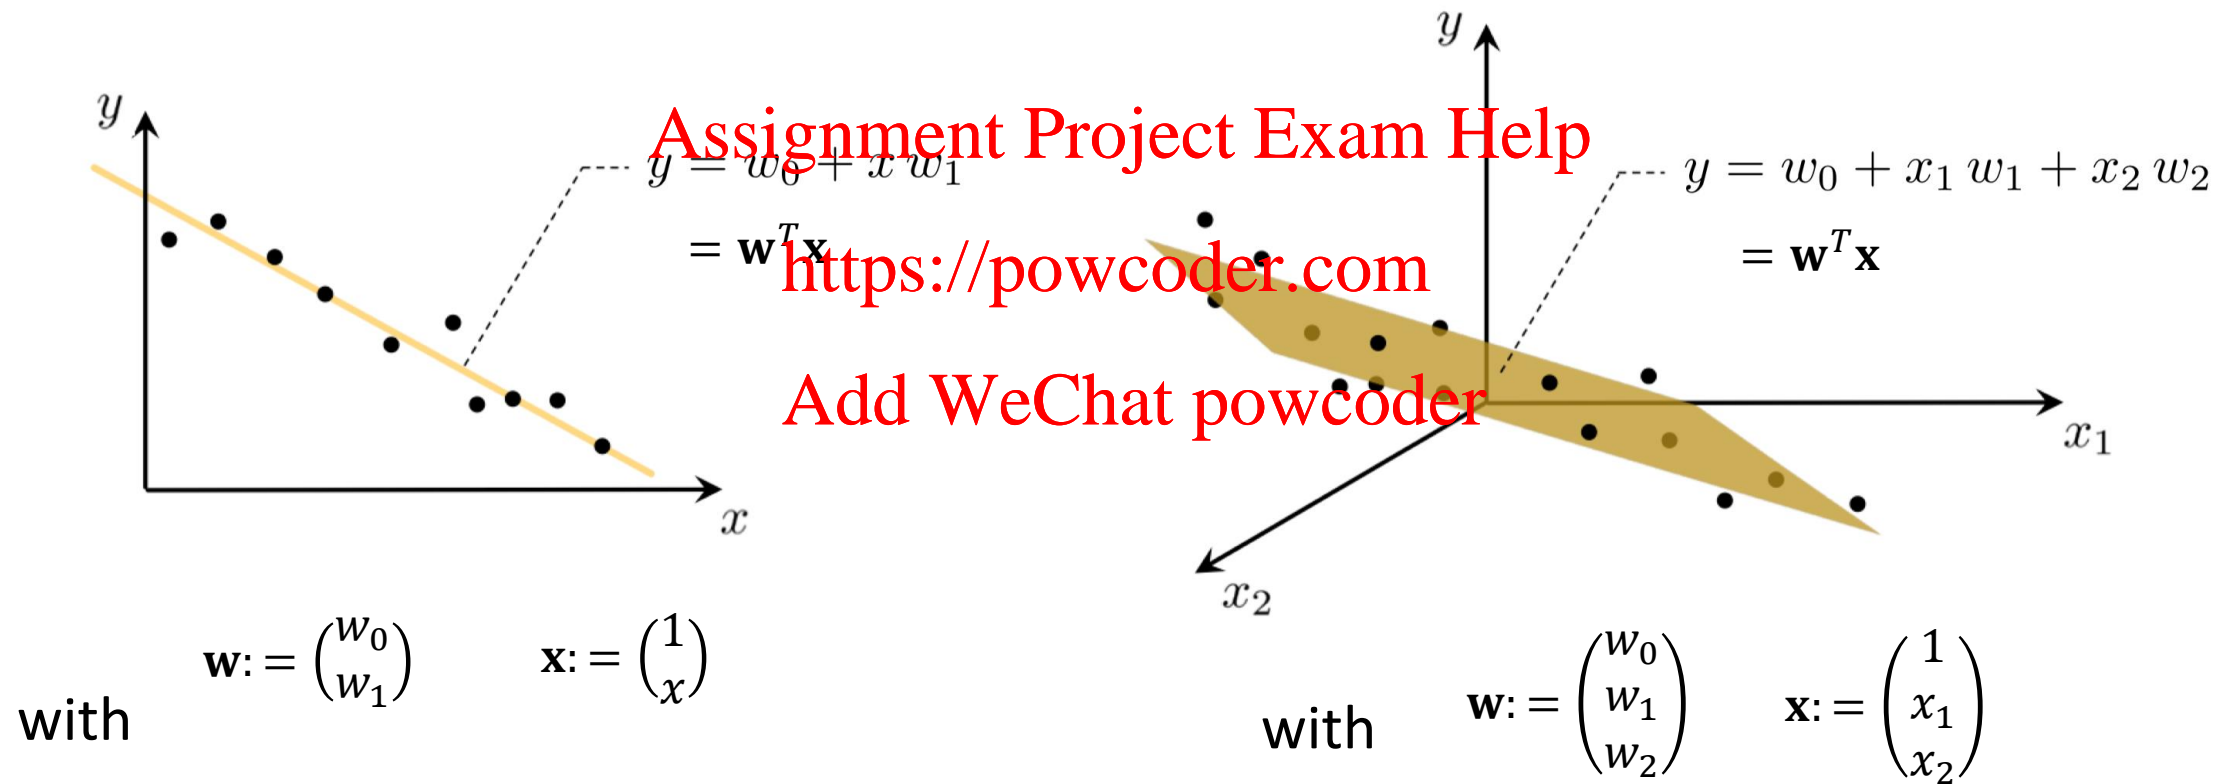
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- We change the model
- The loss and cost function remains the same

Multivariate linear regression



Univariate nonlinear regression

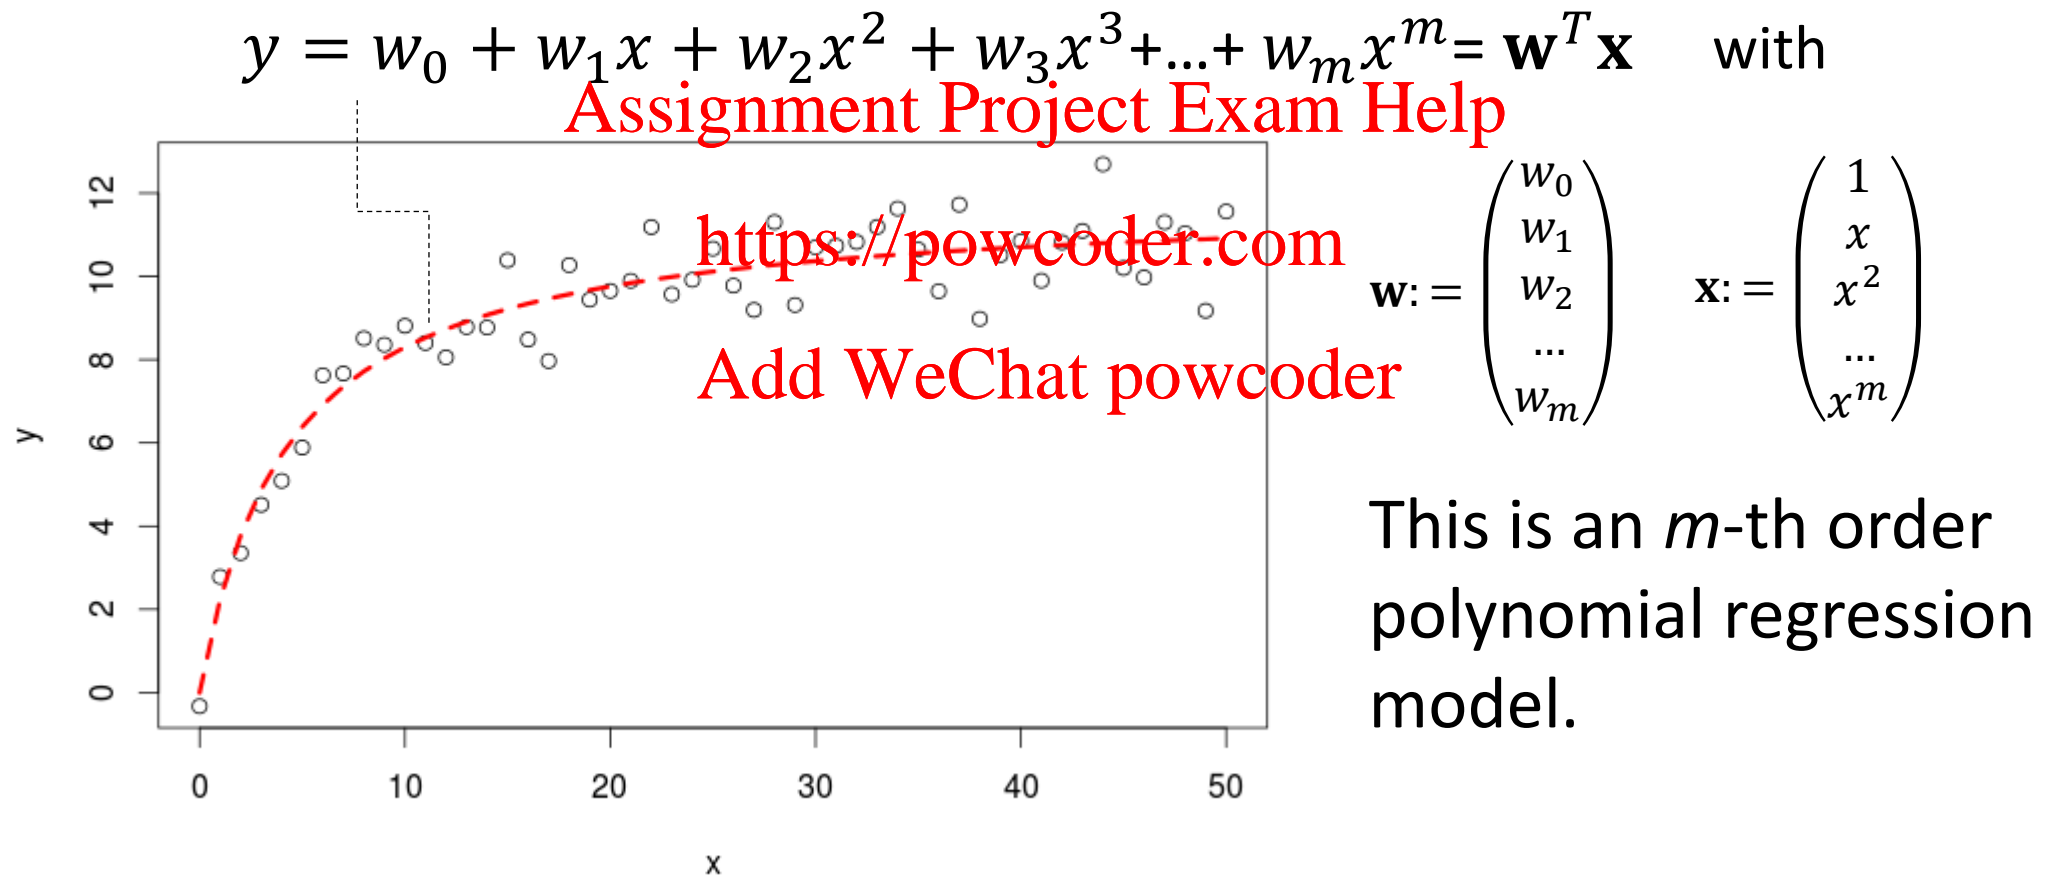



Figure from: <https://www.r-bloggers.com/first-steps-with-non-linear-regression-in-r>

Advantages of vector notation

- Vector notation is concise
- With the vectors \mathbf{w} and \mathbf{x} populated appropriately (and differently in each case, as on the previous 2 slides), these models are still linear in the parameter vector. <https://powcoder.com>
- The cost function is the L2 as before
- So the gradient in both cases is:
$$\nabla g(\mathbf{w}) = 2(\mathbf{w}^T \mathbf{x}^{(n)} - y^{(n)}) \mathbf{x}^{(n)}$$
- Ready to be plugged into the general gradient descent algorithm

Don't get too carried away with nonlinearity



Reference Acknowledgement

Several figures and animations on these slides are taken from:

- Jeremy Watt et al. ~~Machine Learning Refined~~. Cambridge University Press, 2020.

<https://github.com/jeremvatt/machine-learning-refined>

Add WeChat powcoder